

# Receipt Evaluation Report

Github source: <https://github.com/SaKamiAn/FTEC5660>

Colab source: <https://colab.research.google.com/drive/1FTJ8mvW4c3jTVPXHlunY3nVK-5uYVZHe#scrollTo=2D4vHPk2joH>

## 1. Problem Background

In this task, **7 receipt images** are given and we are required to answer two evaluation queries with agentic AI method.

**Query 1:** How much money did I spend in total for these bills?

**Query 2:** How much would I have had to pay without the discount?

The core challenge is that the information (total paid, discounts) must first be **extracted from images using a multimodal LLM (Gemini)**, and then **calculated deterministically by code** to ensure evaluation stability.

## 2. Key Difficulties

### a) Multimodal Output Is Not Guaranteed to Be Pure JSON

Even when explicitly asking for JSON, LLMs may prepend explanations, markdown fences, or natural language text, which causes `json.loads()` to fail.

### b) Separation of Responsibilities

LLMs are good at *information extraction*, but numerical computation must be handled by deterministic code to avoid hallucination and ensure reproducibility.

### c) Multiple Receipts Aggregation

We must first extract per-receipt values, then aggregate across all 7 receipts to answer the queries.

## 3. Solution Strategy

The solution is divided into three clear stages:

### Stage 1: Multimodal Information Extraction

Use the **Google Generative AI SDK (Gemini)** to process 7 receipt images.

Prompt the model to extract, for each receipt:

`total_paid`

`total_discount`

The model is instructed to return results in a strict JSON schema.

### Stage 2: Robust JSON Parsing

Because the model output may still contain extra text, we apply **defensive parsing**:

Extract the first valid JSON object from the model output using a regular expression.

Parse it safely using `json.loads()`.

This step ensures robustness and prevents `JSONDecodeError`.

### **Stage 3: Deterministic Computation**

Once the structured data is available:

**Query 1** is answered by summing all `total_paid` values.

**Query 2** is answered by summing (`total_paid + total_discount`) for each receipt.

All calculations are done purely in Python to ensure correctness and reproducibility.

## **4. Final Output Format**

The final output follows a standardized evaluation-friendly structure:

- a) Per-receipt extracted values
- b) Final numeric answers for Query 1 and Query 2

This design aligns with common benchmark and homework evaluation pipelines.

## **5. Key Takeaways**

LLMs should be treated as **probabilistic extractors**, not trusted calculators.

Always assume LLM output may violate formatting constraints.

Deterministic code is essential for numerical evaluation tasks

## Appendix with some information from Google

I failed with Langchain because of some reports of the error. So, I turned to Google GenAI SDK to meet the requirement of the Homework1. And there are some following reasons.

### 1. Dependency Stability and Environment Compatibility

LangChain introduces multiple transitive dependencies, including optional integrations with deep learning frameworks such as **PyTorch**.

In managed notebook environments (e.g., Google Colab), this may trigger **unexpected runtime conflicts**, such as namespace registration errors in `torch.library`, even when the user does not explicitly invoke any torch-related functionality.

By contrast, the GenAI SDK:

- Has a **minimal dependency surface**
- Does not implicitly load PyTorch or Triton
- Is significantly less prone to environment-level runtime errors

This makes it more suitable for constrained or preconfigured execution environments used in assignments and benchmarks.

### 2. Explicit Control Over Multimodal Inputs

When processing multiple receipt images, precise control over:

- Image encoding (base64)
- Input ordering
- Prompt structure

is critical.

The GenAI SDK allows direct construction of multimodal inputs without additional abstraction layers.

This reduces hidden transformations and makes the data flow **fully explicit and auditable**, which is important for debugging and reproducibility.

### 3. Predictable Output Handling for Evaluation

This task requires **strictly structured outputs** to support deterministic downstream computation and evaluation.

Using the GenAI SDK:

- The raw model output is directly accessible
- Output post-processing (e.g., JSON extraction) is fully controlled by user code
- There is no hidden prompt templating or output wrapping

This aligns well with benchmark-style evaluation pipelines, where **format stability is more important than developer convenience**.

### 4. Clear Separation of Responsibilities

The chosen design enforces a clean separation:

- **LLM** → probabilistic information extraction from images
- **Program code** → deterministic numerical computation and aggregation

Using the GenAI SDK keeps this boundary explicit and avoids over-abstracting LLM calls into chains that may blur responsibility between extraction and reasoning.

### **Summary**

LangChain is a powerful framework for complex agentic workflows.

However, for a **multimodal extraction + deterministic evaluation task**, the GenAI SDK provides:

- Higher runtime stability
- Lower abstraction overhead
- Greater control over inputs and outputs

Therefore, it is the more appropriate choice for this assignment.